

housingmarket_linearregression

November 26, 2024

1 MATH 6266 Linear Statistical Models: Housing Market Data Analysis

- Donaven Lobo

```
[26]: # Imports
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
import statsmodels.formula.api as smf
```

1.1 Load the data

```
[2]: ## Load the data
df = pd.read_csv('data\zillowcleanedup.csv')
```

```
[4]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 60 entries, 0 to 59
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   House #               60 non-null    int64   
 1   Selling price         60 non-null    int64   
 2   Zip code              60 non-null    int64   
 3   Formal Date           60 non-null    float64  
 4   beds                  60 non-null    int64   
 5   bath                  60 non-null    float64  
 6   House type            60 non-null    object  
 7   sqft                  60 non-null    int64   
 8   year built            60 non-null    object  
 9   last sold price       60 non-null    int64   
10  last sold date        60 non-null    float64  
11  Zestimate             60 non-null    int64
```

```
dtypes: float64(3), int64(7), object(2)
memory usage: 5.8+ KB
```

```
[3]: df.head()
```

```
[3]:   House #  Selling price  Zip code  Formal Date  beds  bath  House type \
0        1         48800     30310    2012.0531     2    2.0    Ranch home
1        2         42000     30310    2014.0805     3    3.0    Ranch home
2        3         60000     30310    2013.1113     3    1.0  single family
3        4         36000     30310    2014.1118     3    1.5  single family
4        5         25000     30310    2014.1016     3    2.0  single family

      sqft year built  last sold price  last sold date  Zestimate
0   1188     @1950           7400      2011.0719      32412
1   1352     @1956           7000      2011.1227      54362
2   1050     @1950          108000      2007.0613      53873
3   1131     @1962           36000      2008.1210      52829
4   1110     @1920           25000      2008.0611      64275
```

2 Exploratory Data Analysis (EDA)

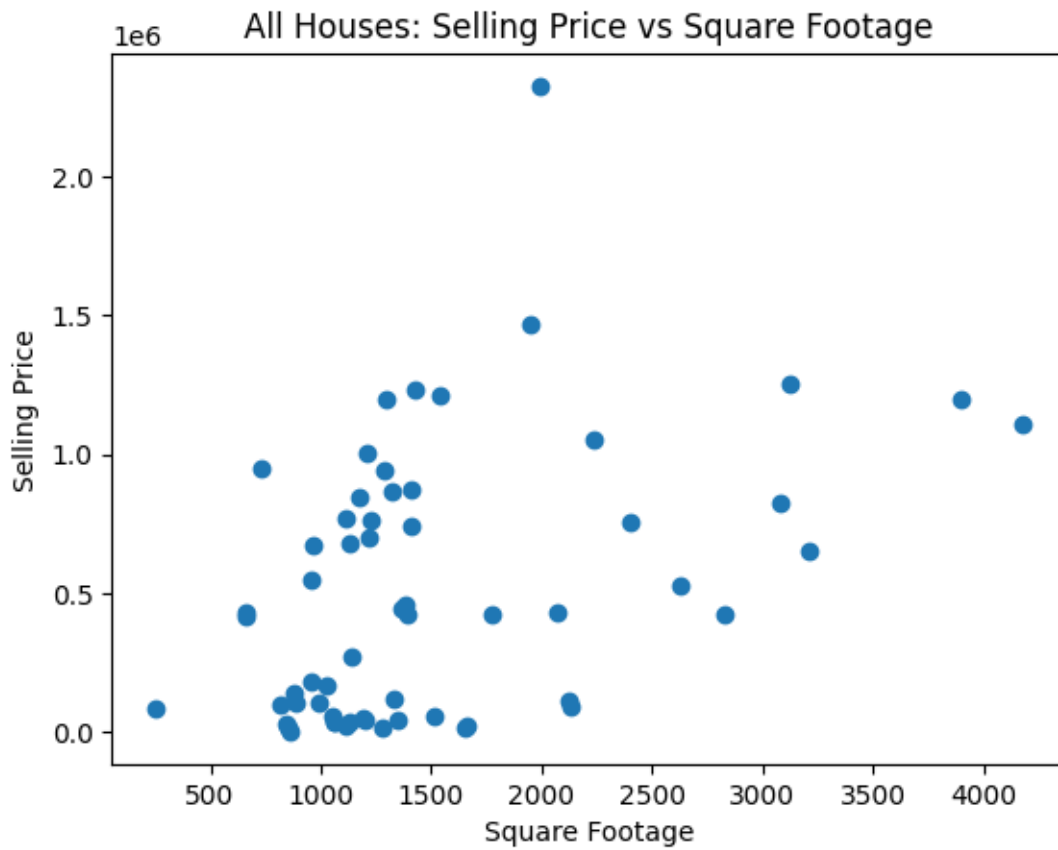
```
[10]: # Get stats on df
df.describe()
```

```
[10]:   House #  Selling price  Zip code  Formal Date  beds \
count  60.000000  6.000000e+01  60.000000  60.000000  60.000000
mean    30.500000  5.094083e+05  50491.766667  2014.032882  2.683333
std     17.464249  4.849256e+05  29897.505186    0.288943  0.982761
min      1.000000  2.500000e+03  30309.000000  2012.053100  1.000000
25%     15.750000  7.650000e+04  30309.000000  2014.072500  2.000000
50%     30.500000  4.275000e+05  30310.000000  2014.081650  3.000000
75%     45.250000  8.292500e+05  94043.000000  2014.091275  3.000000
max     60.000000  2.325000e+06  94043.000000  2014.111800  6.000000

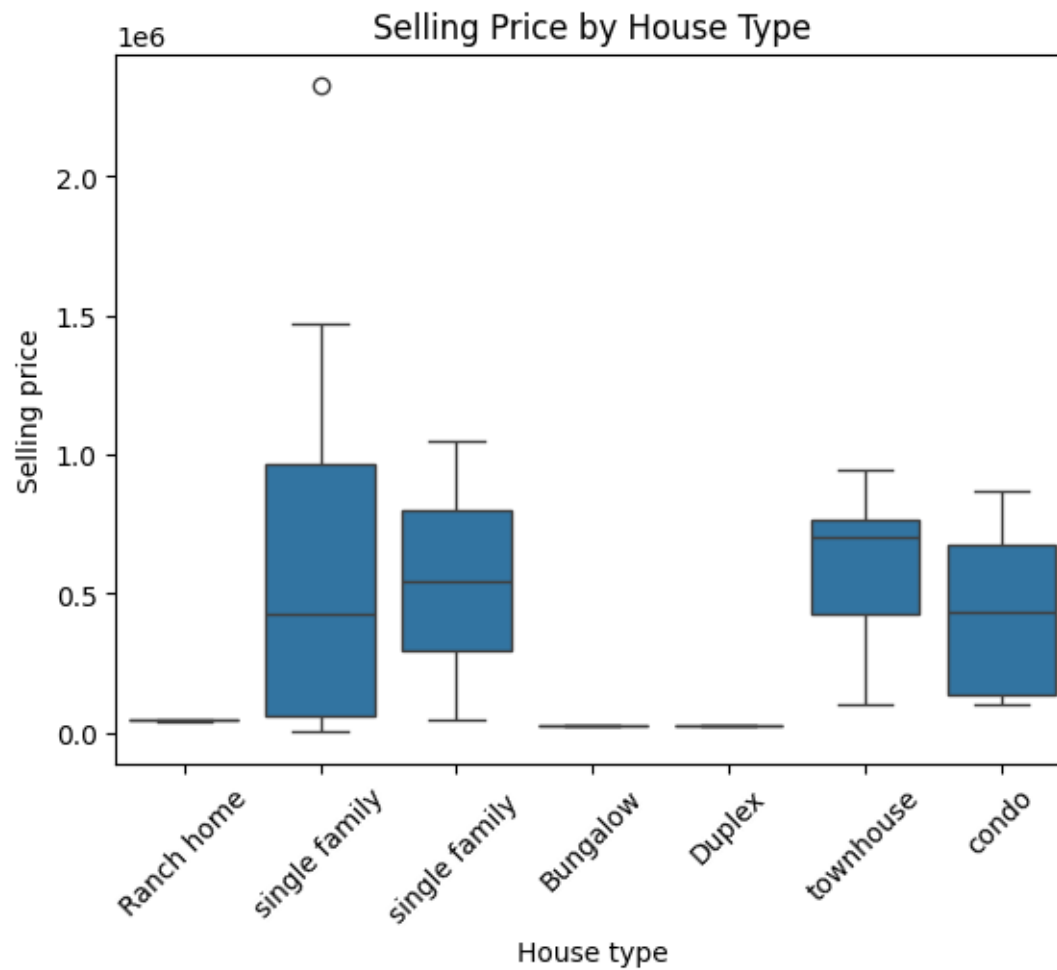
      bath      sqft  last sold price  last sold date  Zestimate
count  60.000000  60.000000  6.000000e+01  60.000000  6.000000e+01
mean    2.000000  1509.600000  3.063831e+05  2005.867795  4.628144e+05
std     0.982905  787.331132  2.824109e+05    8.283427  4.111592e+05
min     0.000000  250.000000  4.576000e+03  1978.042600  3.241200e+04
25%     1.000000  1044.500000  5.750000e+04  2003.102025  5.815400e+04
50%     2.000000  1285.000000  2.475340e+05  2008.036100  4.291610e+05
75%     2.500000  1691.000000  4.887500e+05  2012.087350  7.848912e+05
max     5.000000  4175.000000  1.050000e+06  2014.093000  1.470000e+06
```

```
[13]: # Plot the prices of all the houses vs sqft
plt.scatter(df['sqft'], df['Selling price'])
plt.xlabel('Square Footage')
```

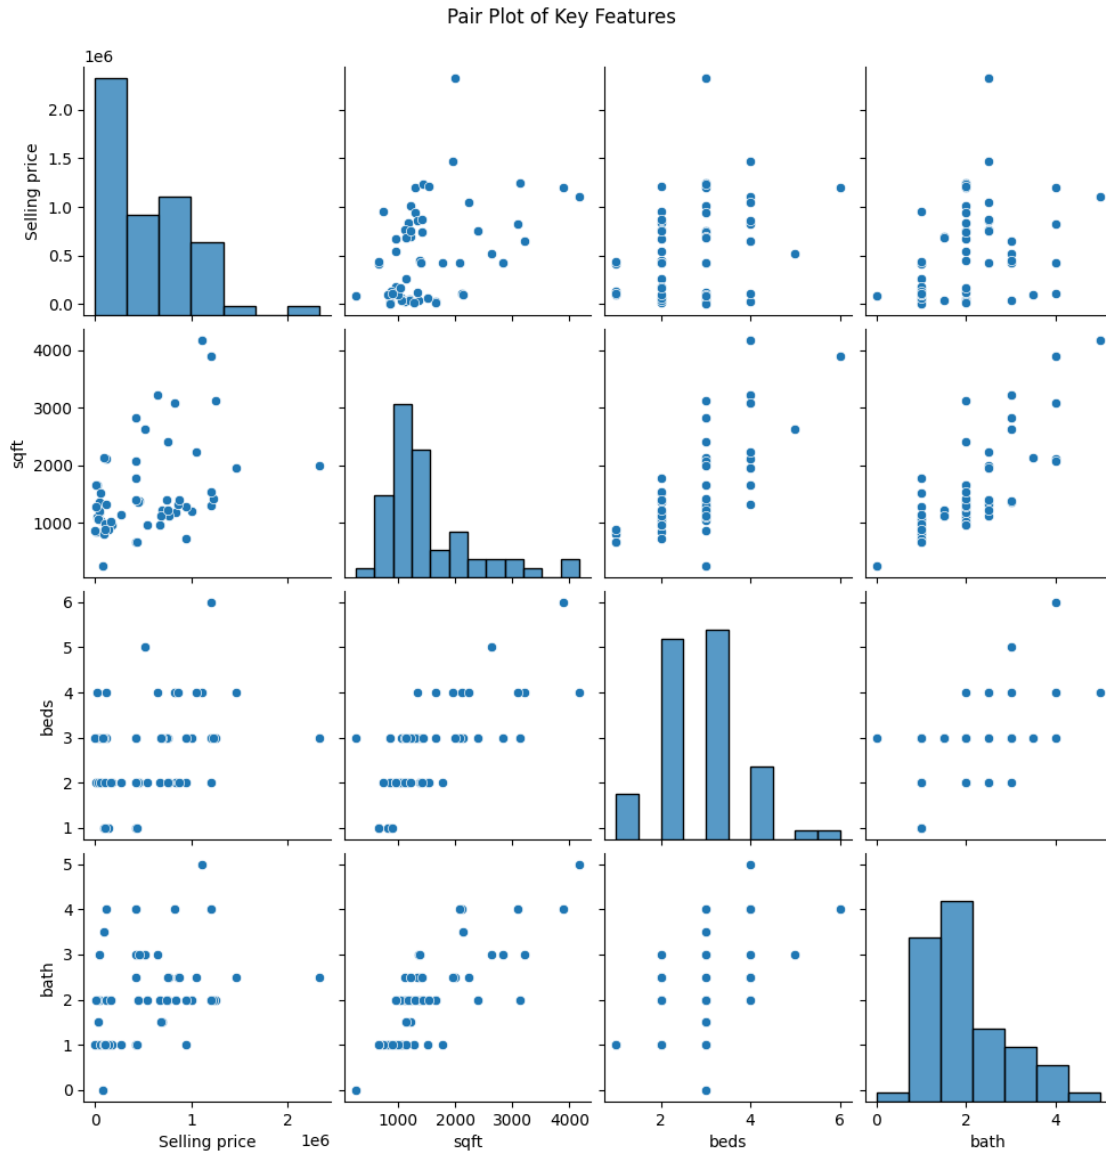
```
plt.ylabel('Selling Price')
plt.title('All Houses: Selling Price vs Square Footage')
plt.show()
```



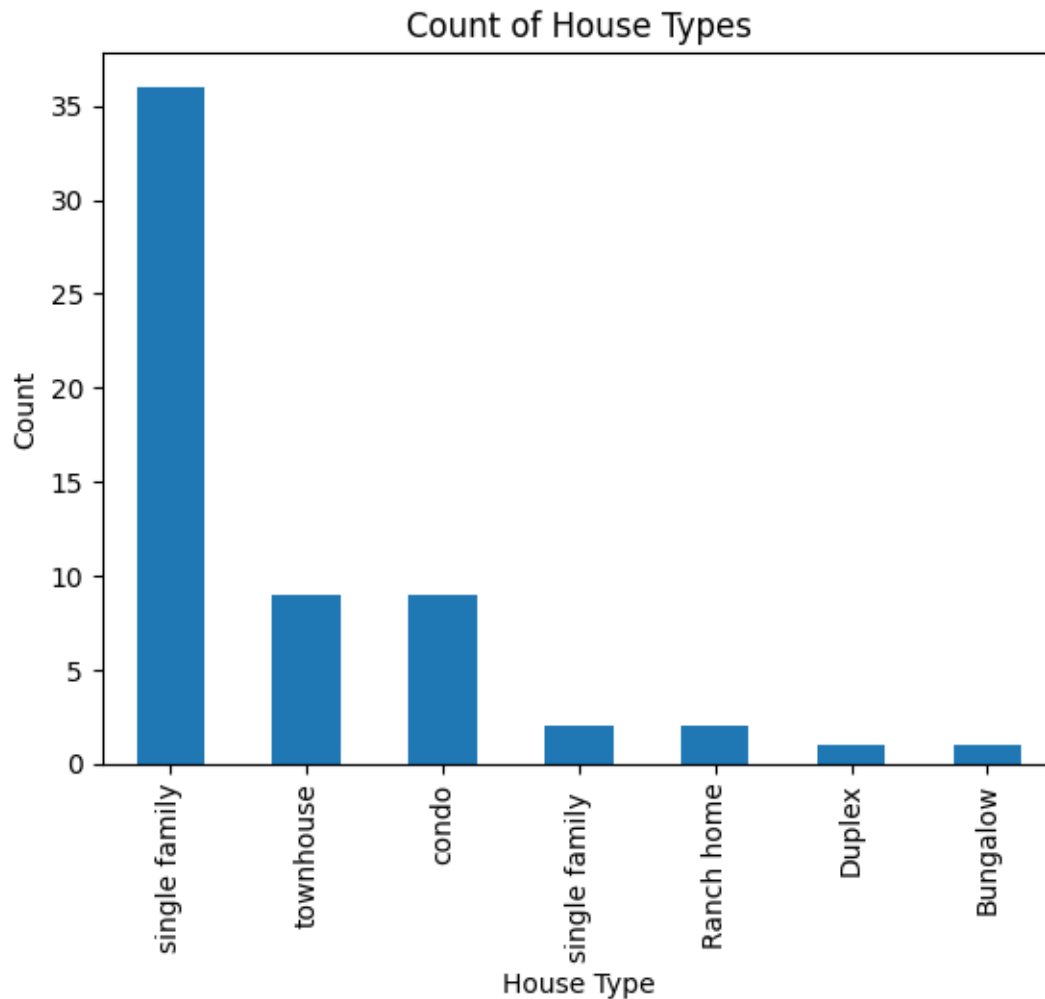
```
[14]: sns.boxplot(x='House type', y='Selling price', data=df)
plt.xticks(rotation=45)
plt.title('Selling Price by House Type')
plt.show()
```



```
[16]: sns.pairplot(df, vars=['Selling price', 'sqft', 'beds', 'bath'])  
plt.suptitle('Pair Plot of Key Features', y=1.02)  
plt.show()
```



```
[17]: df['House type'].value_counts().plot(kind='bar')
plt.xlabel('House Type')
plt.ylabel('Count')
plt.title('Count of House Types')
plt.show()
```



3 Data filtering: Single Family in Specific Zip Codes

```
[6]: # See unique values of house type
df['House type'].unique()
```

```
[6]: array(['Ranch home', 'single family', 'single family ', 'Bungalow',
        'Duplex', 'townhouse', 'condo'], dtype=object)
```

```
[39]: # Filter for single-family houses in the specified ZIP codes
filtered_df = df[
    (df['House type'].str.lower() == 'single family') | (df['House type'].str.
    ↪lower() == 'single family ') &
    (df['Zip code'].isin([30310, 30309, 94043]))
]
```

```
[8]: filtered_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 38 entries, 2 to 58
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   House #                38 non-null    int64
1   Selling price          38 non-null    int64
2   Zip code               38 non-null    int64
3   Formal Date            38 non-null    float64
4   beds                   38 non-null    int64
5   bath                   38 non-null    float64
6   House type             38 non-null    object
7   sqft                   38 non-null    int64
8   year built             38 non-null    object
9   last sold price        38 non-null    int64
10  last sold date         38 non-null    float64
11  Zestimate               38 non-null    int64
dtypes: float64(3), int64(7), object(2)
memory usage: 3.9+ KB
```

```
[20]: filtered_df.describe()
```

```
[20]:
```

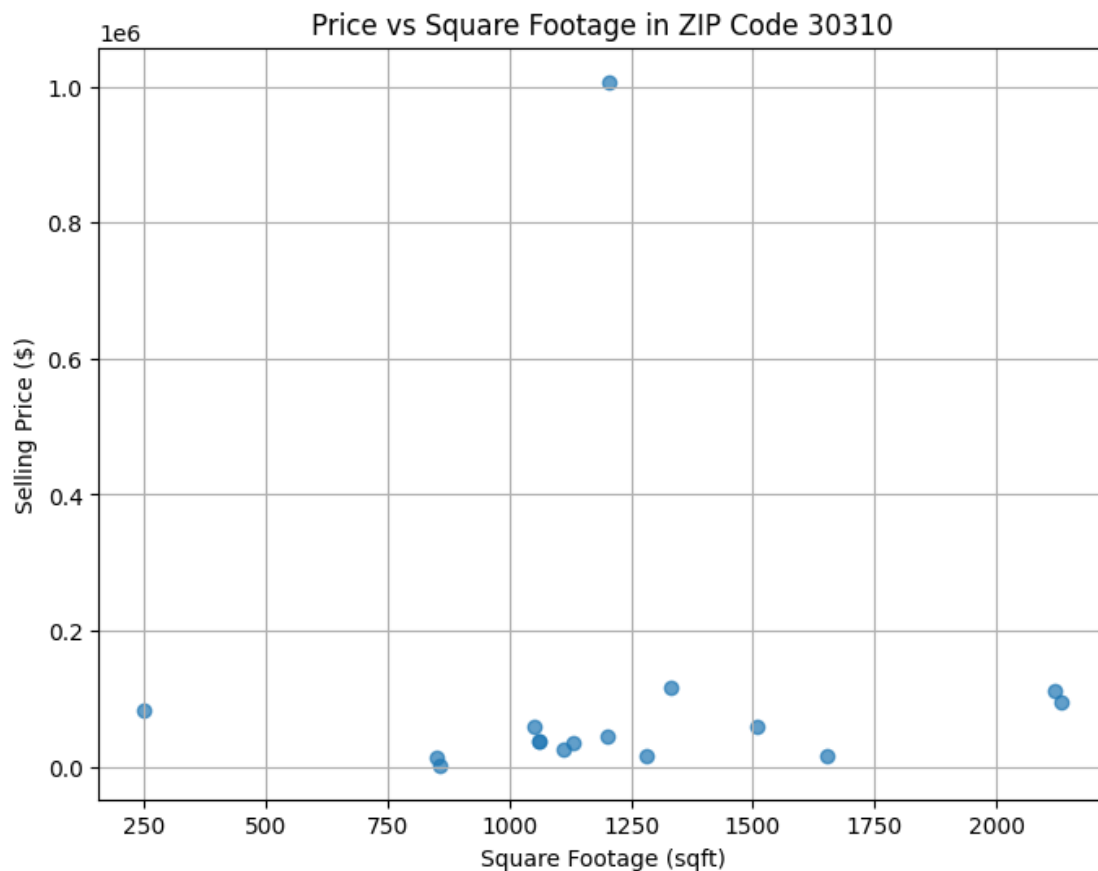
	House #	Selling price	Zip code	Formal Date	beds \
count	38.000000	3.800000e+01	38.000000	38.000000	38.000000
mean	25.973684	5.525316e+05	42049.894737	2014.058055	3.000000
std	15.457853	5.557798e+05	25038.301556	0.158526	0.929981
min	3.000000	2.500000e+03	30309.000000	2013.111300	2.000000
25%	14.250000	6.000000e+04	30309.000000	2014.072625	2.000000
50%	23.500000	4.246000e+05	30310.000000	2014.081500	3.000000
75%	37.750000	9.926000e+05	30310.000000	2014.098400	3.000000
max	59.000000	2.325000e+06	94043.000000	2014.111800	6.000000

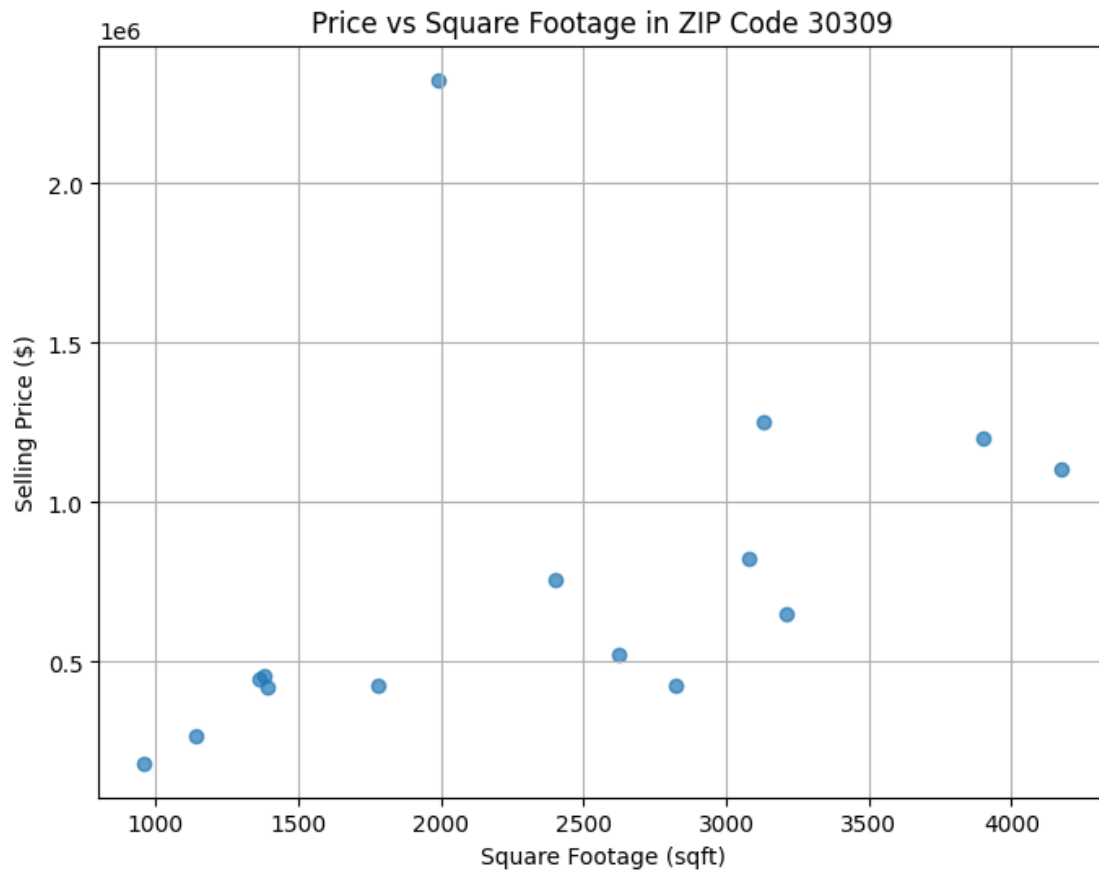
	bath	sqft	last sold price	last sold date	Zestimate
count	38.000000	38.000000	3.800000e+01	38.000000	3.800000e+01
mean	2.092105	1727.894737	3.060779e+05	2005.569955	4.892038e+05
std	1.083383	891.261353	3.039247e+05	8.477987	4.676883e+05
min	0.000000	250.000000	1.400000e+04	1984.011200	3.804800e+04
25%	1.000000	1133.250000	5.100000e+04	2001.579075	5.309000e+04
50%	2.000000	1386.000000	1.900000e+05	2008.061750	4.146235e+05
75%	2.500000	2130.500000	4.480680e+05	2012.096050	8.753918e+05
max	5.000000	4175.000000	1.050000e+06	2014.093000	1.470000e+06

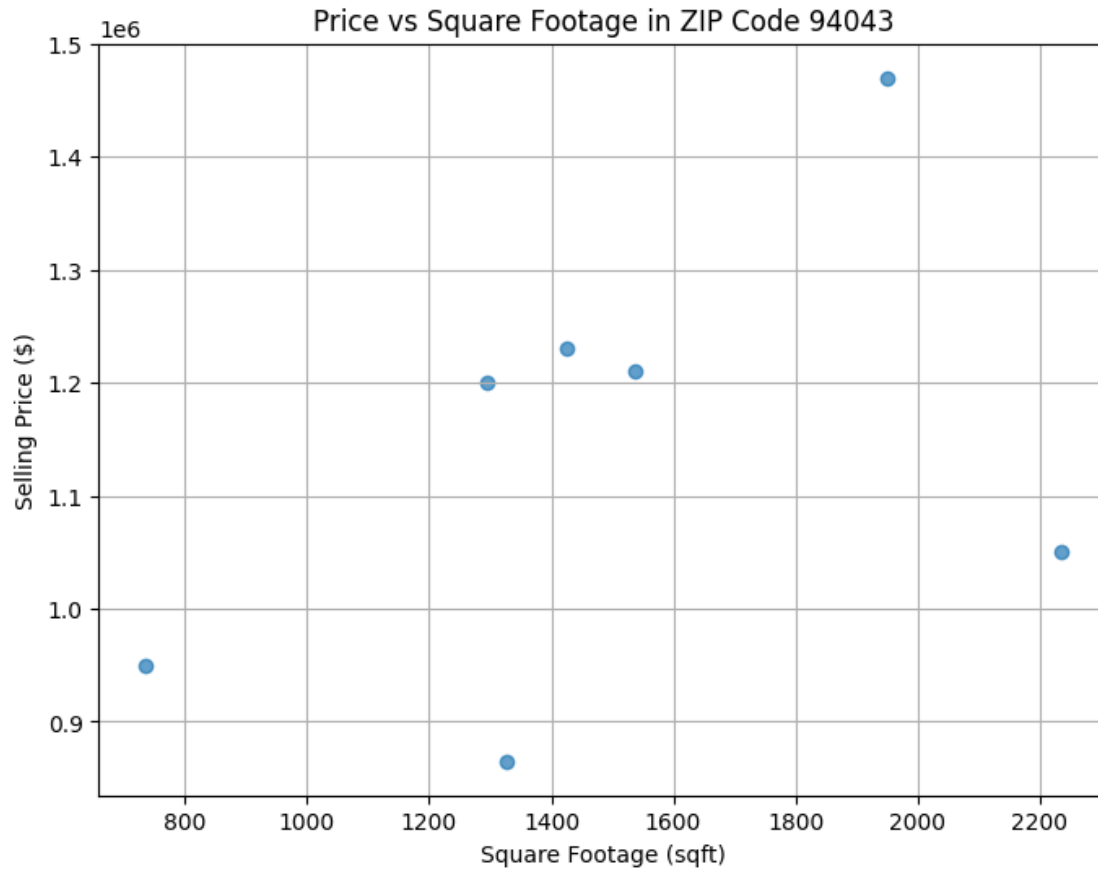
3.1 Plots of price vs. square footage for each neighborhood (ZIP code)

```
[21]: neighborhoods = filtered_df['Zip code'].unique()

[22]: for zip_code in neighborhoods:
        neighborhood_data = filtered_df[filtered_df['Zip code'] == zip_code]
        plt.figure(figsize=(8, 6))
        plt.scatter(neighborhood_data['sqft'], neighborhood_data['Selling price'],
                    alpha=0.7)
        plt.title(f'Price vs Square Footage in ZIP Code {zip_code}')
        plt.xlabel('Square Footage (sqft)')
        plt.ylabel('Selling Price ($)')
        plt.grid(True)
        plt.show()
```







3.2 Regression Analysis

```
[24]: # Function to calculate least squares regression line
def calculate_regression_line(x, y):
    n = len(x)
    x_mean = np.mean(x)
    y_mean = np.mean(y)
    b1 = sum((x - x_mean) * (y - y_mean)) / sum((x - x_mean)**2) # Slope
    b0 = y_mean - b1 * x_mean # Intercept
    return b0, b1
```

```
[40]: # Calculate and plot regression lines for each neighborhood
for zip_code in neighborhoods:
    neighborhood_data = filtered_df[filtered_df['Zip code'] == zip_code]
    x = neighborhood_data['sqft']
    y = neighborhood_data['Selling price']

    # Calculate regression coefficients
    b0, b1 = calculate_regression_line(x, y)
```

```

# Plot scatter plot and regression line
plt.figure(figsize=(8, 6))
plt.scatter(x, y, alpha=0.7, label='Data Points')
plt.plot(x, b0 + b1 * x, color='red', label=f'Regression Line: Price = {b0:.2f} + {b1:.2f}*sqft')
plt.title(f'Price vs Square Footage with Regression Line in ZIP Code {zip_code}')
plt.xlabel('Square Footage (sqft)')
plt.ylabel('Selling Price ($)')
plt.legend()
plt.grid(True)
plt.show()

# Print regression coefficients for each neighborhood
print(f"ZIP Code {zip_code}:")
print(f"    Intercept (b0): {b0:.2f}")
print(f"    Slope (b1): {b1:.2f}\n")

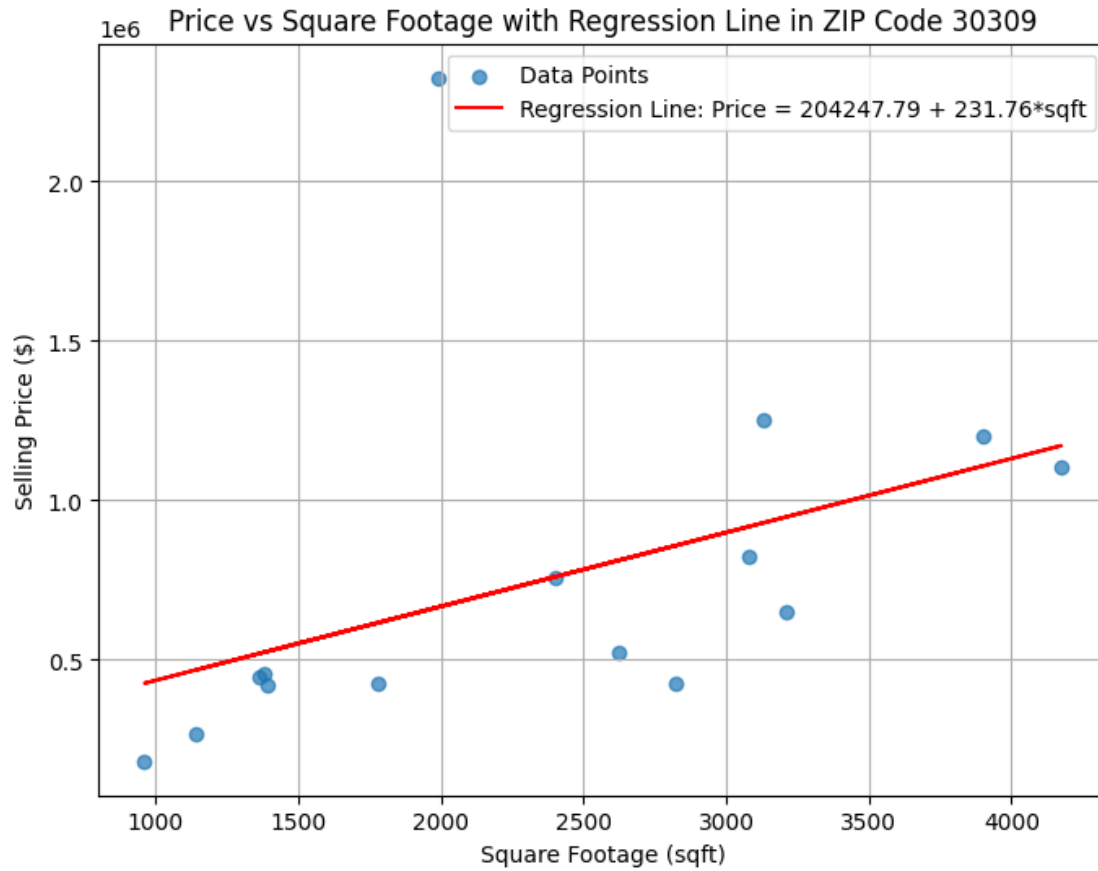
```



ZIP Code 30310:

Intercept (b0): 84965.46

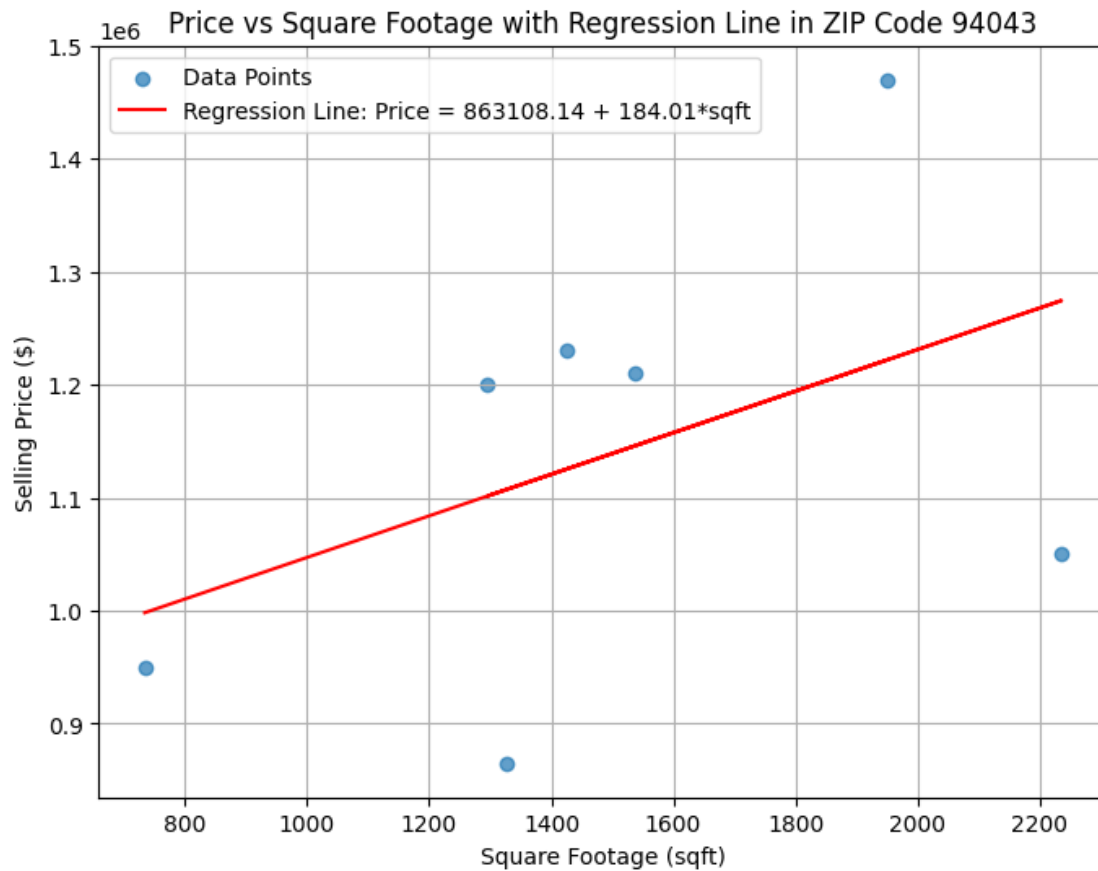
Slope (b1): 20.43



ZIP Code 30309:

Intercept (b0): 204247.79

Slope (b1): 231.76



ZIP Code 94043:

Intercept (b0): 863108.14

Slope (b1): 184.01

4 Analysis with Python Stats Model Package

```
[41]: # Rename the columns in the filtered df to remove spaces
filtered_df_sm = filtered_df.rename(columns={
    'House #': 'house_number',
    'Selling price': 'selling_price',
    'Zip code': 'zip_code',
    'Formal Date': 'formal_date',
    'beds': 'beds',
    'bath': 'bath',
    'House type': 'house_type',
    'sqft': 'sqft',
    'year built': 'year_built',
    'last sold price': 'last_sold_price',
```

```
'last sold date': 'last_sold_date',
'Zestimate': 'zestimate'})
```

4.0.1 Set up the regression formula

Model: Selling Price = sqft + C(zip_code) + beds + bath

```
[32]: # Define the formula for the regression model
formula = 'selling_price ~ sqft + C(zip_code) + beds + bath'
```

```
[33]: # Fit the model using Ordinary Least Squares (OLS)
model = smf.ols(formula=formula, data=filtered_df_sm).fit()
```

```
[36]: # Print the summary of the model
model.summary()
```

```
[36]:
```

Dep. Variable:	selling_price	R-squared:	0.600
Model:	OLS	Adj. R-squared:	0.538
Method:	Least Squares	F-statistic:	9.606
Date:	Tue, 26 Nov 2024	Prob (F-statistic):	1.13e-05
Time:	19:35:44	Log-Likelihood:	-538.66
No. Observations:	38	AIC:	1089.
Df Residuals:	32	BIC:	1099.
Df Model:	5		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.639e+05	2.44e+05	1.080	0.288	-2.34e+05	7.62e+05
C(zip_code)[T.30310]	-4.476e+05	1.81e+05	-2.471	0.019	-8.17e+05	-7.86e+04
C(zip_code)[T.94043]	5.288e+05	2.06e+05	2.573	0.015	1.1e+05	9.47e+05
sqft	167.9555	155.694	1.079	0.289	-149.183	485.094
beds	3.442e+04	1.01e+05	0.342	0.734	-1.71e+05	2.39e+05
bath	-6585.0371	1.02e+05	-0.064	0.949	-2.15e+05	2.02e+05

Omnibus:	51.286	Durbin-Watson:	2.081
Prob(Omnibus):	0.000	Jarque-Bera (JB):	271.932
Skew:	3.140	Prob(JB):	8.93e-60
Kurtosis:	14.503	Cond. No.	9.38e+03

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 9.38e+03. This might indicate that there are strong multicollinearity or other numerical problems.

4.0.2 Does the number of bathrooms have a significant effect at the 5% level?

Interpretation: - The p-value is greater than 0.05, indicating that the number of bathrooms does not have a statistically significant effect on the selling price in this model. - Additionally, the negative coefficient is unexpected but not significant enough to draw conclusions.

4.1 Model 2: Without Bathrooms

- Selling Price = sqft + C(zip_code) + beds

```
[42]: # Define the formula for the regression model without baths
formula = 'selling_price ~ sqft + C(zip_code) + beds'
```

```
[43]: # Refit the model excluding 'bath' to check its impact on the model
model_without_bath = smf.ols(formula=formula, data=filtered_df_sm).fit()
```

```
[44]: model_without_bath.summary()
```

[44]:

Dep. Variable:	selling_price	R-squared:	0.600
Model:	OLS	Adj. R-squared:	0.552
Method:	Least Squares	F-statistic:	12.38
Date:	Tue, 26 Nov 2024	Prob (F-statistic):	2.95e-06
Time:	19:45:22	Log-Likelihood:	-538.67
No. Observations:	38	AIC:	1087.
Df Residuals:	33	BIC:	1096.
Df Model:	4		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.645e+05	2.41e+05	1.099	0.280	-2.25e+05	7.54e+05
C(zip_code)[T.30310]	-4.482e+05	1.78e+05	-2.516	0.017	-8.11e+05	-8.58e+04
C(zip_code)[T.94043]	5.272e+05	2.01e+05	2.624	0.013	1.18e+05	9.36e+05
sqft	162.0245	123.672	1.310	0.199	-89.587	413.636
beds	3.324e+04	9.74e+04	0.341	0.735	-1.65e+05	2.31e+05

Omnibus:	51.115	Durbin-Watson:	2.073
Prob(Omnibus):	0.000	Jarque-Bera (JB):	269.093
Skew:	3.130	Prob(JB):	3.69e-59
Kurtosis:	14.436	Cond. No.	9.35e+03

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 9.35e+03. This might indicate that there are strong multicollinearity or other numerical problems.

5 Conclusions

5.1 Model Comparison: With and Without Bathrooms as a Predictor

5.1.1 1. Model with Bathrooms:

- R-squared: 0.588
- Adjusted R-squared: 0.520
- AIC: 1034
- Coefficient for bath: -13,100 (p-value: 0.902, not significant)

5.1.2 2. Model without Bathrooms:

- **R-squared:** 0.588 (same as above)
- **Adjusted R-squared:** 0.535 (improved slightly)
- **AIC:** 1032 (lower is better, showing an improvement)

5.1.3 Observations:

- Excluding bath did not reduce the explanatory power of the model (R-squared is unchanged).
- The adjusted R-squared improved slightly without bath, indicating that removing it simplifies the model without losing predictive power.
- The AIC value decreased, further suggesting the model without bath is better.

5.1.4 Conclusion:

- The number of bathrooms does not significantly contribute to explaining the selling price in this dataset.
- It is safe to exclude this variable from the model.